Apache Spark and parallel data processing

Pregunta 1:

1.
```
1  rdd = sc.parallelize(range(100))
2  rdd2 = range(100)
3
```

1 punto

Please consider the following code.

Where is the execution of API calls on "rdd" taking place?

- ○ On the local Driver machine
- ◉ In the ApacheSpark worker nodes

Pregunta 2:

2.
```
1  rdd = sc.parallelize(range(100))
2  rdd2 = range(100)
3
```

Please consider the following code.

Where is data in " **rdd2** " stored physically?

- ○ On the local Driver machine
- ◉ In main-memory of ApacheSpark worker nodes

Pregunta 3:

3. What is the parallel version of the following code?

```
1  len(range(9999999999))
2
```

- ◉ sc.parallelize(range(9999999999)).count()
- ○ parallelize(range(9999999999)).count()
- ○ len(sc.parallelize(range(9999999999)))
- ○ size(sc.parallelize(range(9999999999)))
- ○ count(sc.parallelize(range(9999999999)))

Pregunta 4:

4. Which storage solutions support seamless modification of schemas? (Select all that apply)

- [x] ObjectStorage
- [ ] NoSQL
- [x] SQL/Relational Databases

Pregunta 5:

5. Which storage solutions support dynamic scaling on storage? (Select all that apply)

- [x] ObjectStorage
- [ ] NoSQL
- [ ] SQL/Relational Databases

Pregunta 6:

6. Which storage solutions support normalization and integrity checks on data out of the box? (Select all that apply)

- [ ] ObjectStorage
- [ ] NoSQL
- [x] SQL/Relational Databases

Pregunta 7:

7. What is the advantage of using ApacheSparkSQL over RDDs? (select all that apply)

- [ ] ApacheSparkSQL bypasses the RDD interface which has been proven to be very complicated
- [x] SQL is simpler than RDD but has some performance drawbacks
- [x] Catalyst and Tungsten are able to optimise the execution, so are more likely to execute more quickly than if you would had implemented something equivalent using the RDD API.
- [x] The API is simpler and doesn't require specific functional programming skills